

Machine learning

Q-2. Imagine you have a dataset where you have different features like Age , Gender , Height , Weight , BMI , and Blood Pressure and you have to classify the people into different classes like Normal , Overweight , Obesity , Underweight , and Extreme Obesity by using any 4 different classification algorithms. Now you have to build a model which can classify people into different classes. Dataset This is the Dataset You can use this dataset for this question.

In [1]:

```
1  ## Import the necessary libraries:-
2  import pandas as pd
3  from sklearn.model_selection import train_test_split
4  from sklearn.tree import DecisionTreeClassifier
5  from sklearn.linear_model import LogisticRegression
6  from sklearn.ensemble import RandomForestClassifier
7  from sklearn.svm import SVC
8  from sklearn.neighbors import KNeighborsClassifier
9  from sklearn.metrics import classification_report
10 from sklearn.preprocessing import LabelEncoder
11 import warnings
12 warnings.filterwarnings('ignore')
```

In [3]:

```
1  # Load the dataset
2  data = pd.read_csv('Downloads/archive (2)/ObesityDataSet_raw_and_data_sinthetic.csv')
```

In [4]:

```
1  ## Checking top 5 rows
2  data.head()
```

Out[4]:

	Gender	Age	Height	Weight	family_history_with_overweight	FAVC	FCVC	NCP	CA
0	Female	21.0	1.62	64.0	yes	no	2.0	3.0	Someti
1	Female	21.0	1.52	56.0	yes	no	3.0	3.0	Someti
2	Male	23.0	1.80	77.0	yes	no	2.0	3.0	Someti
3	Male	27.0	1.80	87.0	no	no	3.0	3.0	Someti
4	Male	22.0	1.78	89.8	no	no	2.0	1.0	Someti

In [5]:

```
1 ## Checking Rows & Columns Availabale in Dataset
2 data.shape
```

Out[5]:

(2111, 17)

In [6]:

```
1 ## Checking Details Information related with Dataset
2 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Gender                                   2111 non-null   object
1   Age                                     2111 non-null   float64
2   Height                                 2111 non-null   float64
3   Weight                                2111 non-null   float64
4   family_history_with_overweight         2111 non-null   object
5   FAVC                                   2111 non-null   object
6   FCVC                                   2111 non-null   float64
7   NCP                                    2111 non-null   float64
8   CAEC                                   2111 non-null   object
9   SMOKE                                  2111 non-null   object
10  CH2O                                   2111 non-null   float64
11  SCC                                    2111 non-null   object
12  FAF                                    2111 non-null   float64
13  TUE                                    2111 non-null   float64
14  CALC                                   2111 non-null   object
15  MTRANS                                 2111 non-null   object
16  NObeyesdad                             2111 non-null   object
dtypes: float64(8), object(9)
memory usage: 280.5+ KB
```

In [7]:

```
1 ## Checking ALL Columns name present in dataset
2 data.columns
```

Out[7]:

```
Index(['Gender', 'Age', 'Height', 'Weight', 'family_history_with_overweig
ht',
      'FAVC', 'FCVC', 'NCP', 'CAEC', 'SMOKE', 'CH2O', 'SCC', 'FAF', 'TU
E',
      'CALC', 'MTRANS', 'NObeyesdad'],
      dtype='object')
```

In [8]:

```
1 ## Checking Statistical Analysis of Dataset
2 data.describe()
```

Out[8]:

	Age	Height	Weight	FCVC	NCP	CH2O	
count	2111.000000	2111.000000	2111.000000	2111.000000	2111.000000	2111.000000	2111.00
mean	24.312600	1.701677	86.586058	2.419043	2.685628	2.008011	1.01
std	6.345968	0.093305	26.191172	0.533927	0.778039	0.612953	0.85
min	14.000000	1.450000	39.000000	1.000000	1.000000	1.000000	0.00
25%	19.947192	1.630000	65.473343	2.000000	2.658738	1.584812	0.12
50%	22.777890	1.700499	83.000000	2.385502	3.000000	2.000000	1.00
75%	26.000000	1.768464	107.430682	3.000000	3.000000	2.477420	1.66
max	61.000000	1.980000	173.000000	3.000000	4.000000	3.000000	3.00

In [9]:

```
1 ## Checking Information Related with Dataset
2 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender                                2111 non-null   object
1   Age                                   2111 non-null   float64
2   Height                               2111 non-null   float64
3   Weight                               2111 non-null   float64
4   family_history_with_overweight       2111 non-null   object
5   FAVC                                 2111 non-null   object
6   FCVC                                 2111 non-null   float64
7   NCP                                  2111 non-null   float64
8   CAEC                                 2111 non-null   object
9   SMOKE                                2111 non-null   object
10  CH2O                                 2111 non-null   float64
11  SCC                                  2111 non-null   object
12  FAF                                  2111 non-null   float64
13  TUE                                  2111 non-null   float64
14  CALC                                 2111 non-null   object
15  MTRANS                               2111 non-null   object
16  NObeyesdad                           2111 non-null   object
dtypes: float64(8), object(9)
memory usage: 280.5+ KB
```

In [10]:

```
1 ## Checking All Columns Available in dataset
2 data.columns
```

Out[10]:

```
Index(['Gender', 'Age', 'Height', 'Weight', 'family_history_with_overweig
ht',
      'FAVC', 'FCVC', 'NCP', 'CAEC', 'SMOKE', 'CH20', 'SCC', 'FAF', 'TU
E',
      'CALC', 'MTRANS', 'NObeyesdad'],
      dtype='object')
```

In [11]:

```
1 # Preprocess the dataset
2 encoder = LabelEncoder()
3 data['Gender'] = encoder.fit_transform(data['Gender'])
4 data['family_history_with_overweight'] = encoder.fit_transform(data['family_history
5 data['FAVC'] = encoder.fit_transform(data['FAVC'])
6 data['CAEC'] = encoder.fit_transform(data['CAEC'])
7 data['SMOKE'] = encoder.fit_transform(data['SMOKE'])
8 data['SCC'] = encoder.fit_transform(data['SCC'])
9 data['CALC'] = encoder.fit_transform(data['CALC'])
10 data['MTRANS'] = encoder.fit_transform(data['MTRANS'])
11 data['NObeyesdad'] = encoder.fit_transform(data['NObeyesdad'])
```

In [12]:

```
1 ## Checking Details Information related with Dataset
2 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2111 entries, 0 to 2110
Data columns (total 17 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Gender                                     2111 non-null   int32
1   Age                                       2111 non-null   float64
2   Height                                   2111 non-null   float64
3   Weight                                   2111 non-null   float64
4   family_history_with_overweight          2111 non-null   int32
5   FAVC                                     2111 non-null   int32
6   FCVC                                     2111 non-null   float64
7   NCP                                      2111 non-null   float64
8   CAEC                                     2111 non-null   int32
9   SMOKE                                    2111 non-null   int32
10  CH20                                     2111 non-null   float64
11  SCC                                      2111 non-null   int32
12  FAF                                      2111 non-null   float64
13  TUE                                      2111 non-null   float64
14  CALC                                     2111 non-null   int32
15  MTRANS                                   2111 non-null   int32
16  NObeyesdad                              2111 non-null   int32
dtypes: float64(8), int32(9)
memory usage: 206.3 KB
```

In [13]:

```
1 # Split the dataset into features (X) and target (y)
2 X = data.drop('NObeyesdad', axis=1)
3 y = data['NObeyesdad']
```

In [14]:

```
1 # Split the data into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
3
```

In [15]:

```
1 # Decision Tree Classifier
2 dt_clf = DecisionTreeClassifier()
3 dt_clf.fit(X_train, y_train)
4 dt_predictions = dt_clf.predict(X_test)
```

In [16]:

```
1 # Logistic Regression Classifier
2 lr_clf = LogisticRegression()
3 lr_clf.fit(X_train, y_train)
4 lr_predictions = lr_clf.predict(X_test)
```

In [17]:

```
1 # Random Forest Classifier
2 rf_clf = RandomForestClassifier()
3 rf_clf.fit(X_train, y_train)
4 rf_predictions = rf_clf.predict(X_test)
```

In [18]:

```
1 # Support Vector Machine (SVM) Classifier
2 svm_clf = SVC()
3 svm_clf.fit(X_train, y_train)
4 svm_predictions = svm_clf.predict(X_test)
```

In [19]:

```

1 # Print classification reports for each classifier
2 print("Decision Tree Classifier:")
3 print(classification_report(y_test, dt_predictions))

```

Decision Tree Classifier:

	precision	recall	f1-score	support
0	0.92	0.96	0.94	56
1	0.87	0.87	0.87	62
2	0.97	0.92	0.95	78
3	0.93	0.95	0.94	58
4	1.00	1.00	1.00	63
5	0.91	0.89	0.90	56
6	0.94	0.96	0.95	50
accuracy			0.94	423
macro avg	0.93	0.94	0.94	423
weighted avg	0.94	0.94	0.94	423

In [20]:

```

1 print("Logistic Regression Classifier:")
2 print(classification_report(y_test, lr_predictions))

```

Logistic Regression Classifier:

	precision	recall	f1-score	support
0	0.74	0.93	0.83	56
1	0.53	0.42	0.47	62
2	0.58	0.60	0.59	78
3	0.82	0.84	0.83	58
4	0.90	1.00	0.95	63
5	0.54	0.38	0.44	56
6	0.35	0.38	0.37	50
accuracy			0.65	423
macro avg	0.64	0.65	0.64	423
weighted avg	0.64	0.65	0.64	423

In [21]:

```
1 print("Random Forest Classifier:")
2 print(classification_report(y_test, rf_predictions))
```

Random Forest Classifier:

	precision	recall	f1-score	support
0	1.00	0.96	0.98	56
1	0.88	0.94	0.91	62
2	0.99	0.96	0.97	78
3	0.97	0.98	0.97	58
4	1.00	1.00	1.00	63
5	0.89	0.88	0.88	56
6	0.96	0.96	0.96	50
accuracy			0.96	423
macro avg	0.95	0.95	0.95	423
weighted avg	0.96	0.96	0.96	423

In [22]:

```
1 print("SVM Classifier:")
2 print(classification_report(y_test, svm_predictions))
```

SVM Classifier:

	precision	recall	f1-score	support
0	0.71	0.88	0.78	56
1	0.48	0.34	0.40	62
2	0.65	0.33	0.44	78
3	0.77	0.41	0.54	58
4	0.56	1.00	0.72	63
5	0.47	0.48	0.47	56
6	0.43	0.58	0.49	50
accuracy			0.57	423
macro avg	0.58	0.57	0.55	423
weighted avg	0.59	0.57	0.54	423

In []:

```
1
```